

NilStore Network: A Protocol for Decentralized, Verifiable, and Economically Efficient Storage

Contents

NilStore Network: A Protocol for Decentralized, Verifiable, and Economically Efficient Storage

Abstract	3
1. Introduction	3
1.1 Motivation	3
1.2 Key Innovations	3
2. System Architecture	4
2.1 Architectural Layers	4
2.2 The DA Chain (L1)	4
2.3 The Settlement Layer (L2)	4
2.4 The ZK-Bridge	4
3. Data Layer (NilFS)	5
3.1 Object Ingestion and Data Units (DUs)	5
3.2 Erasure Coding and Placement	5
3.3 Autonomous Repair Protocol	6
4. Network Layer (Nil-Mesh)	6
4.1 Heisenberg-Lifted Routing	6
4.2 RTT Attestation and QoS Oracle	7
5. Economic Model (\$STOR-Only)	7
5.1 \$STOR (Staking and Capacity Token)	7
5.2 Fee Market for Bandwidth (\$STOR-1559)	8
5.3 Off-Protocol Payer Adapters	8
6. Consensus and Verification	8
6.1 Overview of the Consensus Model	8
6.2 Proof-of-Useful-Data (PoUD)	9
6.3 Proof-of-Delayed-Encode (PoDE)	9
6.4 Nil-VRF and Epoch Beacon	10
6.5 Proof Coverage and Parameters	11
6.6 Retrieval Receipts and QoS Auditing	11
7. The Deal Lifecycle	13
7.1 Quoting and Negotiation (Off-Chain)	13
7.2 Deal Initiation (On-Chain - L2)	13
7.3 Vesting and Slashing	13
7.4 Multi-Stage Epoch Reconfiguration	14
7.5 L2 Registries & Calls (Normative)	14
7.6 Bandwidth Receipts & BW_root (Normative)	15
8. Advanced Features: Spectral Risk Oracle (σ)	15
9. Governance (NilDAO)	15
9.1 Scope	15
9.x Treasury & Token Distribution Policy (Normative)	16
9.4 Bandwidth Quota, Auto-Top-Up & Sponsors (Normative)	16
9.2 Upgrade Process	16
9.3 Freeze Points	17
10. Roadmap and KPIs	17
10.1 Phased Rollout	17

10.2 Key Performance Indicators (Targets)	17
11. Product UX & Economics (Summary)	17
11.1 Client Flows	17
11.2 Pricing & SP Selection	17
11.3 Redundancy Dial & Auto-Rebalance	18
11.4 Capacity-Aware Entry/Exit	18
11.5 Billing & Spend Guards	18
11.6 Research Isolation	18
Appendix A: Core Cryptographic Primitives and Parameters	19
A.1 Dial Parameters (Baseline Profile “S-512”)	19
A.2 Domain Identifiers and Tags	19
A.3 File Manifest & Crypto Policy (Normative)	19
Appendix B: Threat & Abuse Scenarios and Mitigations	20

NilStore Network: A Protocol for Decentralized, Verifiable, and Economically Efficient Storage

(White Paper v1.1)

Date: 2025-12-01 Authors: NilStore Core Team

Abstract

NilStore is a decentralized storage network designed to provide high-throughput, verifiable data storage with significantly reduced operational overhead. It leverages a novel consensus mechanism based on Proof-of-Useful-Data (PoUD) and Proof-of-Delayed-Encode (PoDE), utilizing a canonical byte representation (plaintext or ciphertext) verified via KZG commitments and timed derivations (Argon2id). By employing a topological data placement strategy (Nil-Lattice), NilStore drastically lowers the hardware barrier to entry, enabling participation from edge devices to data centers. This paper details the system architecture, the NilFS data abstraction layer, the Nil-Mesh routing protocol, the \$STOR-Only economic model, the hybrid L1/L2 settlement architecture, and the underlying cryptographic primitives, including the Nil-VRF epoch beacon, that secure the network.

1. Introduction

1.1 Motivation

While existing decentralized storage protocols have demonstrated the viability of incentive-driven storage, they often rely on computationally intensive Proof-of-Replication (PoRep) stacks requiring significant GPU investment. This centralizes the network around large-scale operators and increases the total cost per byte.

NilStore retains strong cryptographic guarantees while eliminating the “sealing” process, reducing the time to onboard storage to minutes on standard CPUs. This democratization of access increases network resilience through geographic distribution and enables a more efficient storage marketplace.

1.2 Key Innovations

- **Canonical byte-accurate possession:** Storage providers (SPs) keep a canonical byte representation of assigned Data Units (DUs) on disk—either ciphertext or plaintext as selected per deal policy—and prove possession regularly with near-certain full-coverage over time. Deal parameter: `privacy_mode ∈ {"ciphertext" (default), "plaintext"}`. All commitments (C_{root}) and PoDE derivations operate strictly over the chosen canonical bytes.
- **PoUD + PoDE:** PoUD (KZG-based Provable Data Possession over DU canonical bytes) + PoDE (timed window derivations using Argon2id) are the **normative** per-epoch proofs.
- **Nil-Mesh Routing:** Heisenberg-lifted K-shortest paths for optimized latency and Sybil resistance.
- **\$STOR-Only Economy:** A unified economic model using \$STOR for capacity commitment, bandwidth settlement, and governance, explicitly excluding external price oracles.
- **Hybrid Settlement:** A specialized L1 (DA Chain) for efficient proof verification bridged via ZK-Rollup to an EVM L2 for liquidity and composability.

- **Nil-VRF Epoch Beacon:** A BLS12-381-based Verifiable Random Function (VRF) with BATMAN threshold aggregation to ensure unbiased, unpredictable epoch challenges.

2. System Architecture

NilStore employs a hybrid architecture that decouples Data Availability (DA) consensus from economic settlement, optimizing for cryptographic efficiency and ecosystem composability.

2.1 Architectural Layers

1. **Data Layer (NilFS):** Handles object ingestion, erasure coding, and placement.
2. **Network Layer (Nil-Mesh):** Manages peer discovery, routing, and QoS measurement.
3. **Consensus Layer (DA Chain - L1):** Verifies **PoUD (KZG multi-open) + PoDE timing attestations**, manages stake, mints rewards, and generates the epoch beacon (Nil-VRF).
4. **Settlement Layer (L2 Rollup):** Handles economic transactions, liquidity, and governance.

2.2 The DA Chain (L1)

The Data Availability Chain is a minimal L1 (built using Cosmos-SDK/Tendermint BFT) optimized for NilStore's cryptographic operations.

- **Function:** Verifying **KZG openings (multi-open)** for PoUD, enforcing PoDE timing bounds via watcher digests, managing \$STOR staking, executing slashing logic, and running the Nil-VRF beacon. It does not run a general-purpose VM.
- **Required pre-compiles (normative):** (a) **BLAKE2s-256**, (b) **Poseidon** (for Merkle paths), (c) **KZG** (G1/G2 ops; multi-open on BLS12-381), (d) **VDF Verification**, and (e) **BLS Signature Verification/Pairing** (for Nil-VRF).
- **Rationale:** The intensive cryptographic operations required for daily proof verification and beacon generation are best handled natively.

2.3 The Settlement Layer (L2)

Settlement occurs on a ZK-Rollup (using PlonK/Kimchi) bridged to a major EVM ecosystem.

- **Function (One-Token Profile):** Manages **\$STOR only**, mints Deal NFTs, hosts the NilDAO, and executes **\$STOR-denominated** settlement for storage and bandwidth. **Non-\$STOR assets are out-of-scope** for protocol contracts; any conversions happen off-protocol.

2.4 The ZK-Bridge

The L1 aggregates epoch verification results into a single proof/digest and posts it to the L2 bridge contract. **Normative circuit boundary:**

1. **Public inputs:** {epoch_id, DA_state_root, poud_root, bw_root, validator_set_hash}.
2. **Verification key:** vk_hash = sha256(vk_bytes) pinned in the L2 bridge at deployment. Upgrades require DAO action and timelock. An Emergency Circuit MAY perform expedited **VK-only** upgrades under specific conditions (See Section 9.2).
3. **State mapping:** On accept, the bridge **atomically** updates {poud_root, bw_root, epoch_id}; monotonic epoch_id prevents replay.

4. **Failure domains & Yellow-Flag Freeze (Normative):** • On mismatches or non-monotonic epochs, initiate a Grace Period (default 24h). • If a verifier VK emergency is declared by the Emergency Circuit (§9.2), the bridge enters a yellow-flag freeze: continue updating {epoch_id, poud_root, bw_root} but disable all fund-moving paths (vesting, transfers/mints/burns, withdrawals/deposits), new deal creation/uptake, slashing executions, and governance actions except the ratification vote. • Yellow-flag auto-expires after 14 days unless ratified by a full DAO vote. Validity is enforced by the proof and pinned vk_hash; no trusted relayers are required.
5. **Proof Generation (Normative):** The ZK proof MUST be generated by a decentralized prover network or a rotating committee selected from the L1 validator set, with slashing penalties for failure to submit valid proofs.

3. Data Layer (NilFS)

NilFS abstracts data management complexity, automating the preparation, distribution, and maintenance of data. SPs and Users interact with Data Units (DUs), not raw files or replication strategies.

3.1 Object Ingestion and Data Units (DUs)

1. **Content-Defined Chunking (CDC):** Ingested objects are split using CDC (e.g., Rabin fingerprinting) to maximize deduplication. Chunks are organized into a Merkle DAG (CIDv1 compatible).
2. **Data Unit Packing:** Chunks are serialized and packed into standardized **Data Units (DUs)**. DU sizes are powers-of-two (1 MiB to 8 GiB).
3. **Commitment:** The client computes a DU KZG commitment (C_{root}) over the serialized data (See Section 6.2.1). This C_{root} binds the content for all future proofs.

3.2 Erasure Coding and Placement

Normative (Redundancy & Sharding). Each DU is encoded with **systematic Reed-Solomon over GF(2^8)** using **symbol_size = 1 KiB**. The default profile is **RS(n=12, k=9)** ($\approx 1.33\times$ overhead), tolerating $f = 3$ arbitrary shard losses.

3.2.1 Deterministic Placement (Nil-Lattice) Shards are placed on a directed hex-ring lattice (Nil-Lattice) to maximize topological distance and network resilience.

- **Coordinate Calculation:** The coordinate (r, θ) for shard j is determined by: $\text{pos} := \text{Hash}(\text{CID_DU} \parallel \text{ClientSalt_32B} \parallel j) \rightarrow (r, \theta)$
- **Anti-Grinding (Normative):** ClientSalt_32B MUST be derived deterministically from the client's signature over the Deal parameters (e.g., Blake2s-256("NILSTORE-SALT-V1" || Sig_Client)) to prevent placement grinding.
- **Placement constraints (Normative):** At most **one shard per SP per ring-cell**. Shards of the same DU MUST be placed across distinct ring-cells with a minimum topological separation (governance-tunable).

3.2.2 Advanced Profiles (Optional)

- **RS-2D-Hex Profile:** Couples two-dimensional erasure coding with the Nil-Lattice. It maps row redundancy to radial rings and column redundancy to angular slices, enabling $O(|DU|/n)$ repair bandwidth under churn.
- **Mission-Critical (High Redundancy):** Profiles with high shard counts (e.g., RS(30,10)) provide two distinct benefits:
 - **Significant Durability:** The dataset can survive the simultaneous loss of 2/3rds of the storage nodes (e.g., 20 out of 30).
 - **CDN-Grade Throughput:** Retrievals automatically select the fastest subset of nodes (e.g., the fastest 10 out of 30), effectively acting as a high-performance, latency-optimized CDN.
- **Durability Dial Abstraction:** Exposes a user-visible `durability_target` (e.g., 11 nines) that deterministically resolves to a governance-approved redundancy profile (RS-Standard, RS-2D-Hex, or Mission-Critical).

3.3 Autonomous Repair Protocol

The network autonomously maintains durability through a commit-reveal bounty system.

1. **Detection:** If DU availability drops below the resilience threshold (e.g., $k+1$), a `RepairNeeded` event is triggered.
2. **Commit-Reveal:** Repair nodes reconstruct the missing shards and submit solutions in a two-phase process.
3. **Verification and Bounty:** The L1 chain verifies the solution. The Resilience Bounty (default: 5% of the remaining escrowed fee) is awarded to the earliest valid commitment.
 - **Normative (Commitment Drift Prevention):** The repair solution MUST include openings against the original DU KZG commitment (C_{root}). Re-committing a DU is invalid.
4. **Anti-withholding (Normative):** The SP originally assigned the failed shard incurs an immediate penalty strictly greater than the repair bounty (Default: $1.5 \times \text{Bounty}$) and an automatic demerit. The penalty for triggering a repair is significant (Default: 25% of the DU collateral).
5. **Collocation Filter (Normative):** To prevent SPs from repairing their own dropped shards, a filter disqualifies identities that were recently assigned the shard, are within the same IP subnet (/24 IPv4 or /48 IPv6), the same ASN, or exhibit statistically significant RTT Profile Similarity (See Section 4.2).

4. Network Layer (Nil-Mesh)

Nil-Mesh is the network overlay optimized for low-latency, topologically aware routing, utilizing the geometry of the Nil-Lattice.

4.1 Heisenberg-Lifted Routing

- **Mechanism:** Peer IDs (NodeIDs) are mapped (“lifted”) to elements in a 2-step nilpotent Lie group (Heisenberg-like structure) corresponding to their lattice coordinates.
- **Pathfinding:** K-shortest paths ($K=3$) are computed in this covering space and projected back to the physical network. This offers superior latency compared to standard DHTs.

- **Sybil Resistance:** This approach increases Sybil resistance by requiring attackers to control entire topological regions (“Ring Cells”).

4.1.1 Secure Identity Binding (Normative) Peer IDs are securely bound to lattice coordinates (r, θ) through a costly registration process, preventing rapid movement and ensuring that capturing a Ring Cell requires significant capital and time. To register or move, an SP MUST:

1. Bond a minimum amount of \$STOR (Stake_Min_Cell), specific to the target Ring Cell.
2. Compute a Verifiable Delay Function (VDF) proof: $\text{Proof_Bind} = \text{VDF}(\text{NodeID}, r, \theta, \text{difficulty})$.

The DAO MUST periodically update `Stake_Min_Cell` and VDF `difficulty`, raising them automatically if empirical concentration increases.

4.2 RTT Attestation and QoS Oracle

Verifiable Quality of Service (QoS) is crucial for performance, path selection, and preventing Sybil self-dealing (verifying $\text{RTT} > \text{network floor}$).

- **Attestation:** Nodes continuously monitor and sign Round-Trip Time (RTT) attestations with peers.
- **On-Chain Oracle:** A **stake-weighted attester set** posts RTT digests (Poseidon Merkle roots) to the DA chain.

Normative Oracle Procedures:

1. **Challenge-response:** Clients/attesters issue random tokens; SPs must echo tokens within T_{\max} .
2. **VDF Enforcement (Baseline + Conditional Escalation):** Every attestation MUST include a short-delay VDF proof in baseline mode. The VDF input MUST include the random challenge token, computed after receiving the challenge and before transmitting the response, proving the delay occurred within the RTT window. If the anomaly rate exceeds ε_{sys} for ≥ 3 consecutive epochs, increase the VDF delay stepwise until anomalies subside; publish parameters on-chain per epoch. Total cost is capped by the Verification Load Cap.
3. **Diversity & rotation:** The attester set MUST achieve a minimum diversity score (ASN/Region distribution) and assignments are epoch-randomized and committed on-chain.
4. **Slashing:** Equivocation or forged attestations are slashable via on-chain fraud proofs.
5. **Sybil control:** Weight attesters using **quadratic weighting** ($\text{weight} \propto \sqrt{\text{STOR}}$). The total weight of any single entity or correlated group (defined by ASN/Region cluster or RTT Profile Similarity) MUST NOT exceed 20% of the total attester weight.

5. Economic Model (\$STOR-Only)

NilStore employs a unified token economy (\$STOR) to align long-term security incentives with network utility. The protocol explicitly avoids in-protocol stablecoins and external price oracles.

5.1 \$STOR (Staking and Capacity Token)

- **Supply:** Fixed (1 Billion).

- **Functions:** Staking collateral for SPs and Validators; medium of exchange for storage capacity and bandwidth; governance voting power.
- **Sink:** Slashing events.
- **Float Health Monitors (normative):** Publish a **Circulating Float Ratio (CFR)** = (total supply – staked collateral – escrow – DAO/Treasury/vesting – unspent grants)/total supply. Define **yellow/red bands** (default 30% / 25%). Crossing yellow permits a temporary β taper (increase Treasury share) and widening of downward δ within published bands; crossing red triggers the Economic Circuit Breaker (§ 6.6.3). These measures MUST NOT reduce Core security floors (p_{kzg} , R , B_{min}) nor suppress PoUD/PoDE verification.

5.2 Fee Market for Bandwidth (\$STOR-1559)

One-token profile (normative): The protocol uses **\$STOR only** for bandwidth settlement. **No activity-based inflation** is permitted.

Each region r and epoch t defines **BaseFee $_{r,t}$** (in \$STOR per MiB), adjusted EIP-1559-style toward a byte-throughput target U^* . For a payable origin→edge transfer of b bytes:

Burn	= $\beta \cdot \text{BaseFee}_{r,t} \times b$	// burn share in \$STOR
Treasury	= $(1-\beta) \cdot \text{BaseFee}_{r,t} \times b$	// route to Security Treasury
Payout	= $\text{PremiumPerByte} \times b$	// pay provider in \$STOR

Burn-Share Governor (normative). $\beta \in [\beta_{min}, \beta_{max}]$ is DAO-governed (default $\beta=0.95$; bounds [0.90, 1.00]). During a declared Security Escalation (§ 6.6.3), β MAY be temporarily lowered but MUST satisfy $\beta \geq \beta_{emergency_min}$ (default 0.85) and MUST auto-revert after de-escalation. Changes to β and its bounds are time-locked (≥ 24 h).

Update rule (bounded): $\text{BaseFee}_{\{t+1\}} = \text{BaseFee}_{t} \cdot (1 + \delta \cdot (U_t - U^*) / U^*)$ with bounds (default $\pm 12.5\%$). **Operating Bands (normative).** For each region-class, the DAO MUST publish on-chain { U^* band, δ band} and a minimum 72 h timelock for any changes. The controller MUST clamp intra-epoch adjustments to the published band; out-of-band moves require a DAO vote.

Protocol currency invariant: Settlement and escrow contracts **MUST accept \$STOR only**.

5.3 Off-Protocol Payer Adapters

Wallets, edges, and merchant gateways MAY implement **off-protocol adapters** that quote human-readable prices off-chain, acquire *STOR via external venues, and fund payer**STOR escrow***. These adapters are not part of consensus.

6. Consensus and Verification

The economic model is enforced cryptographically through the PoUD+PoDE consensus mechanism on the L1 DA Chain, driven by the Nil-VRF epoch beacon.

6.1 Overview of the Consensus Model

NilStore proofs operate over a **canonical byte representation** selected per deal via **privacy_mode** (default: ciphertext). The objective is to attest, per epoch, that an SP (a) stores the canonical bytes of

their assigned DU intervals (PoUD) and (b) can perform **timed, beacon-salted derivations** over randomly selected windows quickly enough that fetching from elsewhere is infeasible within the proof window (PoDE).

Security anchors: (i) DU **KZG commitment** C_{root} recorded at deal creation. (ii) BLS-VRF epoch beacon for unbiased challenges (Section 6.4). (iii) On-chain **KZG multi-open** pre-compiles. (iv) Watcher-enforced timing digests.

6.2 Proof-of-Useful-Data (PoUD)

PoUD verifies the content correctness of the stored data against the original commitment.

6.2.1 DU Representation & Commitment (Normative) A DU is segmented logically into **1 KiB symbols**. To commit the data using KZG (which operates over the BLS12-381 scalar field), the DU bytes (as selected by `privacy_mode`) MUST be serialized and chunked into 31-byte elements. Each chunk is interpreted as an integer (little-endian) and embedded as a field element. The KZG commitment C_{root} is computed over the polynomial formed by these field elements.

All KZG operations MUST utilize a common, pinned Structured Reference String (SRS) generated via a verifiable Multi-Party Computation (MPC) ceremony (e.g., Perpetual Powers of Tau).

6.2.2 PoUD Mechanism For each epoch and each assigned DU sliver interval:

1. **Challenge Derivation:** The protocol expands the epoch beacon (Section 6.4) to select q distinct, unpredictable symbol indices.
2. **Proof Submission:** The SP MUST provide one or more **KZG multi-open** proofs at the chosen 1 KiB symbol indices proving membership in the DU commitment C_{root} . SPs SHOULD batch using multi-open to minimize calldata.
3. **Verification:** L1 verifies KZG openings using the precompile (Section 2.2).

6.3 Proof-of-Delayed-Encode (PoDE)

PoDE enforces timed locality, ensuring the SP is actively processing the data and not retrieving it on demand. It uses Argon2id for its memory hardness and sequentiality.

6.3.1 The Derive Function (Normative) PoDE relies on the `Derive` function to deterministically compress a canonical DU window into a verifier-recomputable digest, domain-separated by the epoch beacon.

Algorithm (Argon2id Sequential):

```
Derive(canonical_window: bytes, beacon_salt: bytes, row_id: u32, epoch_id: u64, du_id: u128)

tag = "PODE_DERIVE_ARGON_V1"
salt = Blake2s-256(tag || beacon_salt || u32_le(row_id) || u64_le(epoch_id) || u128_le(du_id))
// Hash input to ensure high entropy input to Argon2id
input_digest = Blake2s-256("PODE_INPUT_DIGEST_V1" || canonical_window)
// Parameters (H_t, H_m, H_p) defined by Dial Profile (See Appendix A).
// H_p MUST be strictly 1 to enforce sequentiality (parallelism=1).
```

```

leaf64 = Argon2id(password=input_digest, salt=salt, t_cost=H_t, m_cost=H_m, parallelism=1
Δ_W     = Blake2s-256(canon_window)
return (leaf64, Δ_W)

```

PoDE Recalibration (Normative): The NilDAO MUST periodically recalibrate H_t and H_m parameters to maintain the target Δ_{work} (1s) based on baseline hardware performance.

6.3.2 PoDE Mechanism

1. **Challenge Derivation:** The protocol selects R PoDE windows of size $W = 8\text{ MiB}$ (governance-tunable).
2. **Timed Derivation:** The SP MUST compute $\text{deriv} = \text{Derive}(\text{canon_bytes[interval]}, \text{beacon_salt}, \text{row_id}, \text{epoch_id}, \text{du_id})$ within the proof window using Argon2id parameters (H_t , H_m , $H_p=1$) from the dial profile. The proof includes $H(\text{leaf64}, \Delta_W)$ and the canonical bytes needed for recomputation.
3. **PoDE Linkage (Normative):** The prover MUST include a KZG opening proof π_{kzg} demonstrating that the `canon_window` input bytes correspond exactly to the data committed in C_{root} .
4. **Concurrency & volume (Normative):** The prover MUST satisfy at least R parallel PoDE sub-challenges per proof window (default $R \geq 16$). The aggregate verified bytes per window MUST be $\geq B_{\text{min}}$ (default $B_{\text{min}} \geq 128\text{ MiB}$); only bytes that are both KZG-opened and successfully derived count toward B_{min} .
5. **Verification:** L1 verifies π_{kzg} . Watchers enforce timing via RTT-Oracle and publish pass/fail digests.

6.4 Nil-VRF and Epoch Beacon

NilStore requires unbiased, unpredictable randomness for selecting PoUD/PoDE challenges and the retrieval sampling set. This is provided by the Nil-VRF, a BLS12-381-based Verifiable Random Function (VRF).

6.4.1 Design Choice The Nil-VRF is instantiated as a **BLS-signature-based VRF**: VRF proofs are BLS signatures on `hash_to_G2(msg)`, and verification is a single pairing check. This design is:

- **Uniquely provable:** A single, deterministic proof per (pk, msg) .
- **Deterministically verifiable** on-chain efficiently.
- **Aggregate-friendly:** Supports BATMAN threshold aggregation ($\geq 2/3$ honest).

We follow RFC 9380 for `hash_to_G2` with a NilStore-specific DST: "BLS12381G2_XMD:SHA-256_SSWU_R0_NIL_VRF_H2G".

6.4.2 BATMAN Threshold Aggregation ($t \geq 2/3$) To ensure beacon liveness and security, NilStore uses BATMAN aggregation.

- **Setup:** Committee size N ; threshold $t = \lceil 2N/3 \rceil$. Master key is split using polynomial secret sharing. Participants MUST provide Proof of Possession (PoP) during registration to prevent rogue-key attacks.
- **Per-epoch share posting:** Each participant i publishes their share (pk_i, π_i) where $\pi_i = sk_i \cdot H(\text{epoch_ctr})$.

- **Aggregation:** The aggregator collects t valid shares and computes the aggregate proof using Lagrange interpolation: $\pi_{\text{agg}} = \sum \lambda_i \cdot \pi_i$.

6.4.3 Deterministic Share-Selection (Normative) To prevent the aggregator from grinding the beacon by subset selection:

1. Participants post shares on L1 before the deadline τ_{close} .
2. The aggregator forms $\text{candidate_set} = \{ (\text{pk}_i, \pi_i) \mid \text{posted before } \tau_{\text{close}} \}$ and publishes $\text{candidate_root} = \text{MerkleRoot}(\text{canonical_encode}(\text{pk}_i, \pi_i))$ plus the total candidate count in the epoch transcript.
3. **Grinding Mitigation:** Let $\text{Seed_select} = \text{beacon}_{\{t-1\}}$ (finalized before share posting begins) and $\text{share_id}_i = \text{Blake2s-256}("BATMAN-SHARE" \parallel \text{compress}(\text{pk}_i) \parallel \text{u64_le}(\text{epoch_ctr}))$.
4. **Canonical Set Definition:** Compute $\text{Score}_i = \text{HMAC-SHA256}(\text{Key}=\text{Seed_select}, \text{Message}=\text{share_id}_i)$ over the **entire** candidate_set ; the canonical aggregation set is strictly the t shares with the lowest Score_i values (share_id tie-break). Aggregators MUST attach Merkle proofs for each selected share to candidate_root .
5. Watchers/rollup circuits MUST verify that candidate_root covers every timely share (fraud proofs on omission) and that the selected subset is derived from that full set. The aggregator MUST use this canonical set to compute $(\pi_{\text{agg}}, \text{pk}_{\text{agg}})$.

6.4.4 On-chain Verification & Beacon Derivation The L1 chain verifies the aggregate proof using a single pairing check.

```
require( pairing(pkAgg, H) == pairing(G1_GEN, piAgg) );
y = blake2s256("NIL_VRF_OUT" \parallel compress(pkAgg) \parallel compress(H) \parallel compress(piAgg));
beacon_t = blake2s256("NIL_BEACON" \parallel y);
```

The 32-byte beacon_t feeds PoUD/PoDE challenge derivation (Section 6.2.2, 6.3.2) and seeds the retrieval-sampling RNG (Section 6.6.1).

6.5 Proof Coverage and Parameters

PDP-PLUS Coverage SLO (normative). Define $\text{CoverageTargetDays}$ (default 365). The governance scheduler MUST choose per-epoch index sets (challenge rate q/M) so that for every active DU, the expected fraction of uncovered bytes $(1 - q/M)^T$ after $T=\text{CoverageTargetDays}$ is $\leq 2^{-18}$.

The scheduler MUST be commit-then-sample: indices for epoch t are pseudorandomly derived from the epoch beacon and are not known to SPs before the deadline of epoch $t-1$.

6.6 Retrieval Receipts and QoS Auditing

To account for bandwidth and ensure Quality of Service, clients sign receipts upon successful retrieval, which are then probabilistically sampled for verification.

- **Receipt Schema (Normative):** $\text{Receipt} := \{ \text{CID_DU}, \text{Bytes}, \text{EpochID}, \text{ChallengeNonce}, \text{ExpiresAt}, \text{Tip_BW}, \text{Miner_ID}, \text{payer_id}, \text{Client_Pubkey}, \text{Sig_Ed25519} [, \text{GatewaySig?}, \text{grant_id?}] \}$
- **Eligibility (normative).** Receipts lacking payer_id are ineligible. If grant_id is present, it MUST verify against the payer's "GRANT-TOKEN-V1"

Merkle root. Settlement MUST compute $\text{Burn} = \beta \cdot \text{BaseFee}[\text{region}, \text{epoch}] \times \text{Bytes}$ before Payout, and MUST route $(1-\beta)$ of $\text{BaseFee} \times \text{Bytes}$ to the Security Treasury.

- **Verification model:** Ed25519 signatures are verified off-chain by watchers and/or on the DA chain. The protocol commits to a **Poseidon Merkle root** of receipts (BW_root) and proves byte-sum consistency.
- **Commit requirement (Normative):** For epoch t , SPs MUST have posted $\text{BW_commit} := \text{Blake2s-256}(\text{BW_root})$ by the last block of epoch $t-1$. Failure to post forfeits all bandwidth payouts for t .

6.6.1 Probabilistic Retrieval Sampling A governance-tunable fraction of receipts (default $\geq 5\%$) are verified each epoch.

1. **Sampling Seed (Normative):** Derived from the epoch beacon: $\text{seed}_t := \text{Blake2s-256}(\text{"NilStore-Sample"} \parallel \text{beacon}_t \parallel \text{epoch_id})$
2. **Abuse Score Calculation:** Calculate an Abuse Score $A_{\text{score}}(\text{SP})$ based on historical failures, RTT anomalies, volume spikes, and RTT Profile Similarity.
3. **Risk-Based Sampling (Normative):** The global sampling fraction p is tunable, but the per-SP sampling rate p_{sp} MUST be dynamically adjusted based on $A_{\text{score}}(\text{SP})$.
4. **Honeypot DUs:** DUs created and funded pseudonymously (via a blinded pool) to mimic organic traffic. Any retrieval receipt for a Honeypot DU is automatically selected for 100% verification.

6.6.2 Verification and Enforcement Watchers verify signatures, nonces, RTT transcripts (via QoS Oracle), and inclusion in BW_root .

- **Fail (Minor):** If failures $\leq \varepsilon$ (default 1%), deduct failing receipts and **forfeit all retrieval payouts** for the epoch.
- **Fail (Major):** If failures $> \varepsilon$, forfeit payouts and apply quadratic slashing to bonded \$STOR.

6.6.3 Verification Load Cap and Economic Circuit Breaker (Normative) The total on-chain verification load MUST be capped (DAO-tunable) to prevent DoS. If the cap is reached during a security escalation:

1. **Prioritize High-Risk Receipts:** Sampling prioritizes receipts associated with high abuse scores.
2. **Source Verification Costs:** Excess costs are sourced from the Security Treasury.
3. **Emergency Burn-Share Override + Surcharge (normative):** If the Treasury is insufficient, the protocol MUST temporarily lower β (routing a larger share of BaseFee to the Security Treasury) within $[\beta_{\text{emergency_min}}, \beta]$ and MAY apply a bounded security surcharge σ_{sec} to BaseFee whose revenues are routed 100% to the Security Treasury. Payouts MUST NOT be throttled. Both switches MUST auto-revert after de-escalation or after 14 days (whichever is sooner) and are subject to the standard timelock unless a yellow-flag freeze (§ 2.4) is active.

7. The Deal Lifecycle

7.1 Quoting and Negotiation (Off-Chain)

Clients query Nil-Mesh for SPs near the required lattice slots. SPs respond with Quotes including price, collateral requirements, and QoS caps. The client selects the optimal bundle using the QoS Oracle.

7.2 Deal Initiation (On-Chain - L2)

1. **CreateDeal:** Client calls the function on the L2 settlement contract, posts the Commitment Root (C_{root}), locks the total storage fee in \$STOR escrow, and mints a **Deal NFT** (ERC-721).
2. **MinerUptake:** The selected SP bonds the required \$STOR collateral and commences service.
3. **StorageAttest:** The SP MUST post an attestation tuple $\{sector_id, origin_root, deal_id\}$ committing to the data layout before proofs are counted toward vesting.

7.3 Vesting and Slashing

- **Vesting:** The escrowed fee is released linearly to the SP each epoch, contingent on a valid **PoUD + PoDE** submission.
- **Vesting & Burning Semantics (normative):** Storage escrow is **not** subject to the § 5.2 1559 burn; it vests to the SP on proof of service. Protocol-level burning occurs via (i) the **bandwidth BaseFee** in § 5.2 and (ii) **slashing events** under this § 7.3.1.
- **Consensus Parameters (Normative):**
 - **Epoch Length (T_{epoch}):** 86,400 s (24 h).
 - **Proof Window (Δ_{submit}):** 120 s after epoch end (network scheduling window); DAO-tunable with a normative floor of 60 s to accommodate cross-geo jitter.
 - **Per-replica Work Bound (Δ_{work}):** 1 s (baseline PoDE calibration).

7.3.1 Slashing Rule (Normative) Missed **PoUD + PoDE** proofs trigger a quadratic penalty on the bonded \$STOR collateral, augmented by a correlation factor to penalize clustered failures.

$$\text{Penalty} = \min(0.50, 0.05 \times (\text{Consecutive_Missed_Epochs})^2) \times \text{Correlation_Factor}(F)$$

- **Correlation_Factor(F):**
 - Let $F_{cluster}$ be the fraction of total capacity within a diversity cluster (ASN×region cell, merged with any collocated /24 IPv4, /48 IPv6, or high RTT Profile Similarity) that failed.
 - $\text{Correlation_Factor}(F) = 1 + \alpha \cdot (F_{cluster})^\beta$ (defaults: $\alpha = 1.0$, $\beta = 2.0$), capped at $\text{cap_corr} = 5.0$ with an SP floor $\text{floor_SP} = 1.0$.
 - Governance bounds (time-locked): $\alpha \in [0.5, 2.0]$, $\$[?][2, 4], \text{cap_corr} [3, 8]$. Parameters MUST stay within bounds.
 - If $F_{global} > F^*$ (default 15%), governance MAY cap network-aggregate burn (e.g., 2%/epoch) but MUST NOT waive per-SP penalties. Collocated identities MUST be merged before computing $F_{cluster}$ to prevent Sybil dilution.

7.4 Multi-Stage Epoch Reconfiguration

7.4.0 Objective Ensure uninterrupted availability during committee churn by directing writes to epoch $e+1$ immediately, while reads remain served by epoch e until the new committee reaches readiness quorum.

7.4.1 Metadata Each DU MUST carry `epoch_written`. During handover, gateways/clients route reads by `epoch_written`; if $\text{epoch_written} < \text{current_epoch}$, they MAY continue reading from the old committee until readiness is signaled.

7.4.2 Committee Readiness Signaling New-epoch SPs MUST signal readiness once all assigned slivers are bootstrapped. A signed message: `{epoch_id, SP_ID, slivers_bootstrapped, timestamp, sig_SP}` is posted on L1. When $\geq 2f+1$ SPs signal, the DA chain emits `CommitteeReady(epoch_id)`.

Readiness Audit (Normative). Before counting an SP toward quorum, watchers MUST successfully retrieve and verify a random audit sample of that SP's assigned slivers (sample size $\geq 1\%$ or ≥ 1 sliver, whichever is larger). Failures cause the SP's readiness flag to be cleared and a backoff timer $\Delta_{\text{ready_backoff}}$ (default 30 min) to apply before re-signal.

7.4.3 Routing Rules

- **Writes:** MUST target the current (newest) epoch.
- **Reads:**
 - If $\text{epoch_written} = \text{current_epoch}$, read from current.
 - If $\text{epoch_written} < \text{current_epoch}$, prefer old committee until `CommitteeReady`, then switch to new. Gateways MUST NOT request slivers from SPs that have not signaled readiness.

7.4.4 Failure Modes

- SPs failing to signal by the epoch deadline are slashed per policy.
- If quorum is not reached by $\Delta_{\text{ready_timeout}}$, the DAO MAY trigger emergency repair bounties.
- False readiness is slashable and MAY cause temporary suspension from deal uptake.

7.4.5 Governance Dials DAO-tunable: $\Delta_{\text{ready_timeout}}$ (default 24h), quorum (default $2f+1$), slashing ratios, and the emergency bounty path.

7.5 L2 Registries & Calls (Normative)

- `register_pcv(provider_id, PCV, proof_bundle)` — Provider Capability Vector registry; watcher probes attached and aggregated via BATMAN.
- `register_edge(edge_id, payer_id, max_miss_budget, bond_stor)` — Registers an edge authorized to emit edge-settled receipts on behalf of `payer_id`. `bond_stor` MUST be $\geq f(\max_{\text{miss_budget}})$ (DAO-tunable). Edges are slashable for forged receipts.

- `submit_bw_root(provider_id, epoch, BW_root, served_bytes, med_latency, agg_sig)` — Aggregation of per-chunk receipts into a per-epoch bandwidth root.
- `spawn_hot_replicas(du_id, epoch, Δr, TTL)` — VRF-mediated hot-replica assignment; requires capacity bonds and enforces TTL/hysteresis.

7.6 Bandwidth Receipts & BW_root (Normative)

- **Receipt schema (\$STOR-only):** `{ du_id, chunk_id, bytes, region, t_start, t_end, provider_id, payer_id, [edge_id?, EdgeSig?], [grant_id?], PremiumPerByte, sig_provider }` hashed under "BW-RECEIPT-V1".
- **Aggregation:** leaves → Poseidon Merkle → "BW-ROOT-V1"; providers submit `(provider_id, epoch, BW_root, served_bytes, agg_sig)`.
- **Eligibility (payer-only + A/B):**
 - (A) Edge-settled: `edge_id + EdgeSig` from a payer-registered edge; payable bytes = **origin→edge** only.
 - (B) Grant-token: `grant_id` valid under the payer's "GRANT-TOKEN-V1" Merkle root and unspent. Receipts lacking `payer_id` are ineligible. Settlement MUST compute $\text{Burn} = \beta \cdot \text{BaseFee}[\text{region}, \text{epoch}] \times \text{bytes}$ before Payout and MUST route $(1-\beta)$ of `BaseFee × bytes` to the Security Treasury.

8. Advanced Features: Spectral Risk Oracle (σ)

To manage systemic risk, NilStore incorporates an on-chain volatility oracle (σ).

- **Mechanism:** σ is calculated daily from the Laplacian eigen-drift of the storage demand graph (tracking object-to-region flows), filtered to exclude manipulative patterns (Sybil filtering, high abuse scores). $\sigma_t := ||\Delta\lambda_1 \dots k(\text{Graph}_t)||_2$
- **Application (Dynamic Collateral):** The required collateral for a deal is dynamically adjusted based on internal network volatility (σ). External price volatility is explicitly excluded. $\text{Required_Collateral} := \text{Base_Collateral} \cdot f(\sigma)$
- **Oracle Management (Normative):** $f(\sigma)$ MUST incorporate dampening (e.g., 30-day EMA). The rate of change in Required_Collateral MUST be capped per epoch (e.g., max 10% increase). A grace period (default 72h) is provided for collateral top-ups before liquidation.

9. Governance (NilDAO)

The network is governed by the NilDAO, utilizing stake-weighted (\$STOR) voting on the L2 Settlement Layer.

9.1 Scope

The DAO controls:

- **Economic parameters:** Slashing ratios, bounty percentages, BaseFee adjustment bounds, **burn-share β bounds** ($\beta, \beta_{\min}, \beta_{\max}, \beta_{\text{emergency_min}}$), and **float monitor thresholds/actions** (CFR yellow/red, δ tapering limits, $\sigma_{\text{sec_max}}$).
- **QoS sampling dials:** p (sampling fraction), ε (tolerance), ε_{sys} (system anomaly rate).

- **PoDE/PoUD pressure dials:** R (parallel sub-challenges), B_min (minimum verified bytes), and PoDE calibration (H_t , H_m).
- **Network parameters:** Durability Dial mapping, reconfiguration thresholds, Verification Load Cap.
- **Network upgrades and the treasury.**

9.x Treasury & Token Distribution Policy (Normative)

Scope. This section governs (a) initial token distribution (genesis/IO), (b) Treasury inflows/outflows, and (c) reporting & controls.

Supply & Genesis. Total supply is fixed at 1,000,000,000 \$STOR. The DAO MUST publish a hash-pinned genesis ledger (allocations, vesting schedules, cliffs) and an IO policy, if any, prior to DAO launch (§ 10.1 Phase 2).

Inflows. Treasury inflows MAY include: (i) $(1-\beta) \cdot \text{BaseFee}$ per § 5.2, (ii) earmarked penalties and forfeitures where specified by policy, and (iii) voluntary grants. Slashed collateral not explicitly routed to Treasury is burned.

Outflows. Allowed spend categories: (i) security verification & audits (§ 6.6.3), (ii) protocol grants (grants/credits for bandwidth per § 9.x references), (iii) repair bounties (§ 3.3 top-ups), (iv) client SDK infra. Each category MUST have annual caps and per-tx caps, hash-pinned.

Controls. All outflows require on-chain proposals, a minimum 72 h timelock, and multi-sig execution by distinct roles (Core, Independent Auditor, Community) matching § 9.2 role diversity. The DAO MUST maintain a 12-month security runway before discretionary outflows.

Reporting. A quarterly, hash-pinned Treasury report (CSV ledger + SHA256SUMS) is mandatory.

9.4 Bandwidth Quota, Auto-Top-Up & Sponsors (Normative)

The protocol uses a **hybrid** bandwidth model: each file has an **included quota** (budget reserved per epoch from uploader deposits in \$STOR; verified receipts **debit** \$STOR escrow). On exhaustion, the file enters a **grace tier** with reduced placement weight until **auto-top-up** or **sponsor** budgets restore full weight. APIs: `set_quota`, `set_auto_top_up`, `sponsor`. Governance sets `w_grace`, roll-over caps, region multipliers, price bands, sponsor caps, and ASN/geo abuse discounts.

9.2 Upgrade Process

- **Standard Upgrades:** Require a proposal, a voting period, and a mandatory 72-hour execution timelock.
- **Emergency Circuit (Hot-Patch):** A predefined **5-of-9** threshold can enact **VK-only** emergency patches (e.g., ZK-Bridge VK update).
 - **Key Allocation and Independence (Normative):** The 9 keys are strictly allocated: Core Team (3), Independent Security Auditor (3 distinct entities), Community/Validator Rep (3). The 5-of-9 threshold MUST include at least one valid signature from each group.
 - **Sunset Clause (Normative):** Emergency patches automatically expire 14 days after activation unless ratified by a full DAO vote. The emergency patch mechanism MUST NOT be capable of modifying the Sunset Clause duration.

9.3 Freeze Points

The cryptographic specification and the tokenomics parameters are hash-pinned and frozen prior to external audits and the formal DAO launch.

10. Roadmap and KPIs

10.1 Phased Rollout

1. **MVP SDK (Rust/TS):** (Completed 2025-09)
2. **DAO Launch & Tokenomics Freeze:** (2025-11)
3. **Public Testnet-0 (L1 DA Chain):** (2026-01) - PoUD+PoDE, Nil-VRF, basic economics.
4. **Edge-Swarm Beta (Retrieval Economy):** (2026-04) - Mobile client, \$BW activated, QoS Oracle.
5. **Rollup Bridge Mainnet (L2 Settlement):** (2026-06) - EVM L2 integration, ZK-Bridge, Deal NFTs.
6. **Mainnet-1:** (2026-09).

10.2 Key Performance Indicators (Targets)

Metric	Target
Onboarding Time (Plaintext)	Minutes (CPU-bound)
Epoch Proof Size (Aggregated)	$\leq 1.2 \text{ kB}$ (post-recursion)
Retrieval RTT (p95)	$\leq 400 \text{ ms}$ (across 5 geo regions)
On-chain Verify Gas (L2 Bridge)	$\leq 120\text{k Gas}$
Durability	$\geq 11 \text{ nines}$ (modeled)
Sampling FP/FN rate	$\leq 0.5\% / \leq 0.1\%$ (monthly audit)
VRF Beacon Verification Gas (L1)	$\approx 97 \text{ k Gas}$ (1 pairing)

11. Product UX & Economics (Summary)

11.1 Client Flows

- `store(file, durability_target, term)` returns `{deal_id, profile, price_quote, placement_summary, estimated_retrieval_price}`; SDKs expose escrow balance, spend caps, redundancy status, and alerts.
- `get(file_id)` auto-selects healthy SPs, shows a retrieval price quote, and enforces a user-set `max_monthly_spend` unless overridden.

11.2 Pricing & SP Selection

- Providers post **bounded price curves** per `{region, QoS}` `{p0, k, γ, cap_free_GiB, min_term, price_curve_id}` within caps/bounds (defaults: $\beta_{\text{floor}}=0.70, \beta_{\text{ceiling}}=1.30, \text{premium}_{\text{max}}=0.5 \times \text{BaseFee}, \text{price}_{\text{cap}}_{\text{GiB}}=2 \times$ the regional/class median BaseFee, $\sigma_{\text{sec}}_{\text{max}} \leq 10\%/\text{epoch}$; k, γ bounded per PSet). The chain publishes a single AskBookRoot + partition offsets each epoch; off-book SPs are not eligible for PoUD/PoDE payouts.

- Deterministic assignment uses {CID_DU, ClientSalt, shard_index, region, qos} to sort candidate slices by marginal price (current util), then QoS, then sp_id, enforcing placement (one shard per SP per ring-cell with min ring/slice distance). Quotes outside caps are rejected; deals must prove inclusion against AskBookRoot.

11.3 Redundancy Dial & Auto-Rebalance

- Durability presets map to pinned profiles: Standard=RS(12,9), Archive=RS(16,12), Mission-Critical=RS-2D-Hex{rows=4, cols=7}. Deals record durability_target and resolved profile.
- Placement constraints per profile: Standard ring distance ≥ 2 ; Archive ring distance ≥ 3 ; Mission-Critical ring distance ≥ 3 and slice distance ≥ 2 ; one shard per SP per cell.
- Auto-rebalance restores target redundancy when degraded or when an SP exits; repairs must open against the original C_root and respect placement diversity. Defaults: T_repair_max = 24h (RS), T_repair_max = 8h (RS-2D-Hex).

11.4 Capacity-Aware Entry/Exit

- Entry probation ramps rewards 50→100% over N_probation = 7 epochs with a slashing multiplier $\lambda_{\text{entry}} = 1.25$.
- Exit fee and unbonding window scale with capacity headroom: headroom_raw = free_capacity_ratio / target_headroom with default target_headroom = 0.20 (pinned to preserve ~20% spare for repairs/unbonding). Optionally smooth headroom_raw via 7-day EMA; use headroom = clamp(0,1, headroom_raw_smoothed). F_exit = F_base × (1 + k_fee × (1 - headroom)) with defaults F_base=0.5%, k_fee=2.0, bounds [0.5%, 10%]; T_unbond = T_base + k_time × (1 - headroom) with defaults T_base=24h, k_time=72h, bounds [12h, 7d]. Surplus capacity saturates to headroom=1 → cheaper/faster exits; tight capacity → costlier/slower exits with mandatory handoff/repair before finalization.

11.5 Billing & Spend Guards

- Single escrow per deal (storage + baseline egress) in \$STOR; auto top-up optional; grace mode reduces QoS weight and pauses new replicas when under-funded. Defaults: K_epoch=7 epochs funded; K_low=3 epochs trigger grace. Retrieval billing is per epoch with bounded β and PremiumPerByte within [0, premium_max] and price_cap_GiB.
- Events (DealCreated, RedundancyDegraded, RepairScheduled, RepairComplete, ProofMissed, ExitRequested/Finalized, FreezeActivated/Cleared, SpendGuardHit) are emitted for UX and observability. Yellow-flag freeze pauses withdrawals/new deals/exits but keeps proofs and billing running; auto top-ups pause.

11.6 Research Isolation

- PoS²-L RFCs are research-only; production profiles remain PoUD+PoDE on canonical bytes unless a DAO-ratified, auto-sunset research activation is explicitly in force.

Appendix A: Core Cryptographic Primitives and Parameters

This appendix details the normative cryptographic primitives, parameters, and policies used in NilStore (Core v2.0).

A.1 Dial Parameters (Baseline Profile “S-512”)

A **dial profile** defines the core cryptographic parameters and the Proof-of-Delayed-Encode (PoDE) settings.

Symbol	Description	Baseline “S-512”
Curve	Elliptic Curve (for KZG and VRF)	BLS12-381 (Mandatory)
r	BLS12-381 subgroup order	0x73EDA753299D7D483339D80809A1D8055
H_t	PoDE Argon2id time cost (iterations)	3 (Calibrated for $\Delta_{\text{work}}=1\text{s}$)
H_m	PoDE Argon2id memory cost (KiB)	1048576 (1 GiB)
H_p	PoDE Argon2id parallelism	1 (Mandatory Sequential)

A.2 Domain Identifiers and Tags

DomainID : u16 partitions digests by purpose. Digests are computed as: digest = Blake2s-256(Version || DomainID || payload)

Core domain strings used with Blake2s-256 across modules:

Tag	Purpose
"NIL_VRF_OUT"	VRF output compression
"BLS12381G2_XMD:SHA-256_SSBU_R0_NIL_VRF_H2G"	VRF hash_to_G2 DST
"NIL_BEACON"	Epoch beacon derivation from VRF output
"NilStore-Sample"	Retrieval-sampling seed from epoch beacon
"PODE_DERIVE_ARGON_V1"	PoDE Derive function tag
"PODE_INPUT_DIGEST_V1"	PoDE input data hashing
"BATMAN-SHARE"	Deterministic share-selection label
"NILSTORE-SALT-V1"	Placement Anti-Grinding Salt derivation

A.3 File Manifest & Crypto Policy (Normative)

NilStore uses a content-addressed file manifest and a deterministic encryption policy.

- **Root CID** = Blake2s-256("FILE-MANIFEST-V1" || CanonicalCBOR(manifest)).
- **File Master Key (FMK)** (32B) is HPKE-wrapped to authorized retrieval keys ("FMK-WRAP-HPKE-V1").

- **Deterministic Key/Nonce Derivation (Normative):** Per-DU Content Encryption Keys (CEKs) and Nonces are derived deterministically from the FMK. (`CEK_32B`, `Nonce_12B` = `HKDF-SHA256(IKM=FMK, info="DU-KEYS-V1" || du_id, L=44)`).
- **AEAD: AES-256-GCM** using the derived CEK and the deterministic 96-bit Nonce.
- **Security Warning (Nonce Reuse):** This deterministic derivation is secure ONLY under the strict assumption that DUs are immutable (Write-Once) and that `du_id` is unique for every distinct plaintext under the same FMK.
- **DU CID** = `Blake2s-256("DU-CID-V1" || ciphertext || tag)`.

Appendix B: Threat & Abuse Scenarios and Mitigations

Scenario	Attack surface	Detect / Prevent (Design)	Normative anchor(s)
Wash-retrieval / Self-dealing	SP scripts fake clients to inflate bandwidth usage	Risk-based sampling (Abuse Score); Challenge-nonce + expiry in receipts; watchers verify Ed25519; Poseidon receipt root commitment; RTT Oracle verification (> network floor).	§6.6 (Receipts/Sampling), §4.2 (QoS Oracle)
RTT Oracle collusion	Gateways/attesters collude to post low RTT	Stake-weighted attesters (quadratic weighting, influence cap); challenge-response tokens; mandatory VDF proof in RTT window; ASN/region diversity; randomized assignments; slashable fraud proofs.	§4.2 (RTT Oracle)
Commitment drift in repair	Repaired shards bound to a <i>new</i> commitment	Repaired shards MUST open against the original DU KZG commitment (<code>C_root</code>); reject new commitments.	§3.3 (Autonomous Repair)

Scenario	Attack surface	Detect / Prevent (Design)	Normative anchor(s)
Beacon Grinding	Aggregator selects subset of VRF shares to bias beacon	Deterministic Share Selection: Canonical set defined by lowest scores derived from previous epoch's beacon. Aggregator cannot choose the input set.	§6.4.3 (Nil-VRF)
Bridge/rollup trust	VK swap or replay of old epoch	L2 bridge pins vk_hash; monotone epoch_id; timelocked VK upgrades; 5-of-9 Emergency Circuit with strict role diversity and sunset clause.	§2.4 (ZK-Bridge), §9.2 (Governance)
Lattice capture (Sybil/cartel)	SPs concentrate shards topologically	One-shard-per-SP-per-cell; minimum cell distance; Secure Identity Binding (Stake_Min_Cell + VDF proof required to move/register coordinates).	§3.2.1 (Placement), §4.1.1 (Identity Binding)
Shard withholding (availability)	SP stores but doesn't serve (or drops data)	Vesting tied to valid PoUD + PoDE; slashing for missed epochs (quadratic + correlation factor); Anti-withholding penalties (immediate penalty > repair bounty).	§7.3 (Vesting/Slashing), §3.3 (Anti-withholding)
PoDE Pre-computation	SP computes derivations before the challenge window	Derive function input is salted with the epoch beacon (unpredictable); Argon2id enforces sequentiality ($H_p=1$); Watchers enforce timing digests within Δ_{submit} .	§6.3 (PoDE), §6.4 (Nil-VRF)